

```
[cloudera@quickstart ~]$ hive
```

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties

WARNING: Hive CLI is deprecated and migration to Beeline is recommended.

```
hive> create database flight;
```

```
OK
```

Time taken: 10.032 seconds

```
hive> use flight;
```

```
OK
```

```
hive> create table flight_det(
```

```
  > fno int,
```

```
  > source varchar(20),
```

```
  > destination varchar(20));
```

```
OK
```

Time taken: 1.971 seconds

```
hive> alter table flight_det add columns (year int);
```

```
OK
```

Time taken: 0.972 seconds

```
hive> describe flight_det;
```

```
OK
```

```
fno                int
```

```
source             varchar(20)
```

```
destination        varchar(20)
```

```
year               int
```

Time taken: 1.059 seconds, Fetched: 4 row(s)

```
hive> alter table flight_det rename to flight_details;
```

```
OK
```

Time taken: 0.484 seconds

```
hive> describe flight_details;
```

```
OK
```

```
fno                int
```

```
source             varchar(20)
```

```
destination        varchar(20)
```

```
year               int
```

Time taken: 0.265 seconds, Fetched: 4 row(s)

```
hive> create external table flight_info(
```

```
  > year INT,
```

```
  > month INT,
```

```
  > day INT,
```

```
  > day_of_week INT,
```

```
  > dep_time INT,
```

```
  > crs_dep_time INT,
```

```
  > arr_time INT,
```

```
  > crs_arr_time INT,
```

```
  > unique_carrier STRING,
```

```
  > flight_num INT,
```

```
  > tail_num STRING,
```

```
  > actual_elapsed_time INT,
```

```
  > crs_elapsed_time INT,
```

```
  > air_time INT,
```

```
  > arr_delay INT,
```

```
  > dep_delay INT,
```

```
  > origin STRING,
```

```
  > dest STRING,
```

```
  > distance INT,
```

```
  > taxi_in INT,
```

```
  > taxi_out INT,
```

```

> cancelled INT,
> cancellation_code STRING,
> diverted INT,
> carrier_delay STRING,
> weather_delay STRING,
> nas_delay STRING,
> security_delay STRING,
> late_aircraft_delay STRING)
> row format delimited fields terminated by ','
> location "/home/cloudera/Desktop/flight_info.csv";
OK
Time taken: 0.155 seconds
hive>

hive> load data local inpath "/home/cloudera/Desktop/flight_info.csv" overwrite into table flight_info;
Loading data to table mydatabase.flight1
Table mydatabase.flight1 st_deats: [numFiles=1, numRows=0, totalSize=25113900, rawDataSize=0]
OK
Time taken: 1.196 seconds

hive> select * from flight_info;
2008      1      1      2      820      820      1238      1254      US      1      N656AW198      214
          177     -16     OPHX     ORD      1440      4       17       0          0       NA      NA
          NA      NA      NA
2008      1      1      2      1013     950      1456      1429      US      3      N653AW223      219
          178     27     23     PHX     ORD      1440      7       38       0          0       0
          0       7       0       20
2008      1      1      2      1502     1320     1936      1750      US      5      N667AW214      210
          187     106    102     PHX     ORD      1440      6       21       0          0       0
          0       106     0       0
2008      1      1      2      1845     1602     2305      2025      US      7      N628AW200      203
          179     160    163     PHX     ORD      1440      8       13       0          0       160
          0       0       0       0
2008      1      1      2      1945     1947      20       19       US      9      N679AW215      212
          174      1      -2     PHX     ORD      1440      27      14       0          0       NA
          NA      NA      NA      NA
2008      1      1      2      2258     2300     313       324       US      606     N654AW195      204
          170     -11     -2     PHX     ORD      1440      12      13       0          0       NA
          NA      NA      NA      NA
2008      1      1      2      829      829      1012      1033      US      86      N315AW163      184
          141     -21     OPHX     PDX      1009      3       19       0          0       NA      NA
          NA      NA      NA
2008      1      1      2      1106     1058     1242      1256      US      88      N314AW156      178
          141     -14     8PHX     PDX      1009      2       13       0          0       NA      NA
          NA      NA      NA
2008      1      1      2      1502     1457     1640      1653      US      90      N673AW158      176
          132     -13     5PHX     PDX      1009      2       24       0          0       NA      NA
          NA      NA      NA
2008      1      1      2      1757     1755     1932      1950      US      92      N627AW155      175
          133     -18     2PHX     PDX      1009      3       19       0          0       NA      NA
          NA      NA      NA

```

2008	1 133 NA	1 -23 NA	2 2PHX NA	2105 PDX	2103 1009	2234 3	2257 13	US 0	94	N837AW149 0 NA	174 NA
2008	1 230 NA	1 -30 NA	2 -2 NA	1008 PHX NA	1010 PHL	1626 2075	1656 4	US 24	250 0	N676AW258 0	286 NA
2008	1 237 NA	1 -13 NA	2 -1 NA	2309 PHX NA	2310 PHL	527 2075	540 6	US 15	258 0	N833AW258 0	270 NA
2008	1 274 NA	1 224 NA	2 -34 NA	817 2PHX NA	815 PHL	1415 2075	1449 3	US 11	700 0	N191UW 0	238 NA
2008	1 228 NA	1 1 NA	2 1PHX NA	1116 PHL	1115 2075	1756 24	1755 28	US 0	1192	N171US 280 0 NA	280 NA
2008	1 224 NA	1 -18 NA	2 6PHX NA	1612 PHL	1606 2075	2216 3	2234 17	US 0	1194	N170US 244 0 NA	268 NA
2008	1 226 0	1 41 0	2 67 0	1407 PHX 41	1300 PHL	2012 2075	1931 6	US 13	1530 0	N173US 245 0	271 0
2008	1 206 NA	1 -23 NA	2 2PHX NA	952 PIT	950 1813	1537 5	1600 14	US 0	1574	N753US 225 0 NA	250 NA
2008	1 208 NA	1 0 NA	2 9PHX NA	1619 PIT	1610 1813	2203 4	2203 12	US 0	1584	N111US 224 0 NA	233 NA
2008	1 208 NA	1 12 NA	2 48 NA	1053 PHX NA	1005 RDU	1637 1891	1625 5	US 11	634 0	N815AW224 0	260 NA
2008	1 87 NA	1 -4 NA	2 4PHX NA	1106 RNO	1102 601	1155 5	1159 17	US 0	522	N605AW109 0 NA	117 NA
2008	1 85 NA	1 -14 NA	2 -3 NA	1425 PHX NA	1428 RNO	1509 601	1523 5	US 14	524 0	N158AW104 0	115 NA
2008	1 78 NA	1 -28 NA	2 -2 NA	2110 PHX NA	2112 RNO	2143 601	2211 4	US 11	528 0	N802AW93 0	119 NA
2008	1 47 NA	1 -19 NA	2 1PHX NA	825 SAN	824 304	831 3	850 16	US 0	180	N901AW66 0 NA	86 NA
2008	1 52 0	1 32 0	2 33 0	2140 PHX 19	2107 SAN	2159 304	2127 5	US 22	188 0	N649AW79 0	80 13

2008	1 47 NA	1 0 NA	2 OPHX NA	1950 SAN	1950 304	2016 2	2016 37	US 0	191	N303AW86 0 NA	86 NA
2008	1 58 NA	1 12 NA	2 12 NA	2305 PHX NA	2253 SAN	2322 304	2310 3	US 16	192 0	N633AW77 0	77 NA
2008	1 47 0	1 18 0	2 40 0	1530 PHX 7	1450 SAN	1531 304	1513 3	US 11	196 0	N302AW61 0	83 11
2008	1 79 NA	1 48 NA	2 5 NA	1115 10 NA	1105 PHX NA	1129 SAN	1124 304	US 5	1427 21	N751UW 0	74 0
2008	1 54 0	1 45 0	2 54 0	1854 PHX 37	1800 SAN	1906 304	1821 2	US 16	1509 0	N167US 72 0	81 8
2008	1 81 NA	1 48 NA	2 -14 NA	2118 3PHX NA	2115 SAN	2122 304	2136 3	US 13	1511 0	N192UW 0	64 NA
NULL	NULL NULL NULL	NULL NULL NULL	NULL NULL NULL	NULL NULL NULL	NULL NULL NULL	NULL NULL NULL	NULL NULL NULL	NULL NULL NULL	NULL NULL NULL	NULL NULL NULL	NULL NULL NULL

Time taken: 10.12 seconds, Fetched: 257250 row(s)

hive>show tables;

OK

flight_details

flight_info

Time taken: 0.392 seconds, Fetched: 2 row(s)

hive> create table air_flight(

> fno int,
> pno int,
> pname varchar(20),
> class varchar(20));

OK

Time taken: 0.941 seconds

hive> insert into air_flight values(101,20,"Parth","Business class");

Query ID = cloudera_20240423052828_92ea2c1a-e723-4dd2-94be-9149135a1640

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1713870747829_0001, Tracking URL =

http://quickstart.cloudera:8088/proxy/application_1713870747829_0001/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1713870747829_0001

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2024-04-23 05:35:48,168 Stage-1 map = 0%, reduce = 0%

2024-04-23 05:36:49,724 Stage-1 map = 0%, reduce = 0%

2024-04-23 05:37:50,501 Stage-1 map = 0%, reduce = 0%

2024-04-23 05:38:51,567 Stage-1 map = 0%, reduce = 0%

2024-04-23 05:39:43,804 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 8.2 sec

MapReduce Total cumulative CPU time: 8 seconds 200 msec

Ended Job = job_1713870747829_0001

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/flight.db/air_flight/.hive-staging_hive_2024-04-23_05-28-01_806_7474073666891730511-1/-ext-10000

Loading data to table flight.air_flight

Table flight.air_flight stats: [numFiles=1, numRows=1, totalSize=28, rawDataSize=27]

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 9.0 sec HDFS Read: 4135 HDFS Write: 101 SUCCESS

Total MapReduce CPU Time Spent: 9 seconds 0 msec

OK

Time taken: 724.805 seconds

hive> insert into air_flight values(102,21,"Ayush","Economy class");

Query ID = cloudera_20240423054040_32b568ee-d59d-44ba-ac33-bae369e5472b

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1713870747829_0002, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1713870747829_0002/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1713870747829_0002

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2024-04-23 05:43:35,455 Stage-1 map = 0%, reduce = 0%

2024-04-23 05:44:35,681 Stage-1 map = 0%, reduce = 0%

2024-04-23 05:45:36,872 Stage-1 map = 0%, reduce = 0%

2024-04-23 05:46:13,293 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 7.06 sec

MapReduce Total cumulative CPU time: 7 seconds 60 msec

Ended Job = job_1713870747829_0002

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/flight.db/air_flight/.hive-staging_hive_2024-04-23_05-40-19_185_5887912077937198594-1/-ext-10000

Loading data to table flight.air_flight

Table flight.air_flight stats: [numFiles=2, numRows=2, totalSize=55, rawDataSize=53]

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 7.73 sec HDFS Read: 4224 HDFS Write: 100 SUCCESS

Total MapReduce CPU Time Spent: 7 seconds 730 msec

OK

Time taken: 368.089 seconds

hive> select * from air_flight;

OK

101	20	Parth	Business class
-----	----	-------	----------------

102	21	Ayush	Economy class
-----	----	-------	---------------

Time taken: 0.64 seconds, Fetched: 2 row(s)

hive> describe flight_details;

OK

fno	int
-----	-----

source	varchar(20)
--------	-------------

destination	varchar(20)
-------------	-------------

year	int
------	-----

Time taken: 1.418 seconds, Fetched: 4 row(s)

hive> insert into flight_details values(101,"Pune","Mumbai",2022);

Query ID = cloudera_20240423055252_02fa91b5-7c04-4533-a465-e10933174f68

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1713870747829_0003, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1713870747829_0003/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1713870747829_0003

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2024-04-23 05:56:31,517 Stage-1 map = 0%, reduce = 0%

2024-04-23 05:57:32,359 Stage-1 map = 0%, reduce = 0%

2024-04-23 05:58:33,277 Stage-1 map = 0%, reduce = 0%

2024-04-23 05:59:34,289 Stage-1 map = 0%, reduce = 0%

2024-04-23 06:00:34,489 Stage-1 map = 0%, reduce = 0%

2024-04-23 06:01:34,807 Stage-1 map = 0%, reduce = 0%

2024-04-23 06:02:36,926 Stage-1 map = 0%, reduce = 0%

2024-04-23 06:03:39,318 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 5.67 sec

2024-04-23 06:03:59,070 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 10.8 sec

MapReduce Total cumulative CPU time: 10 seconds 800 msec

Ended Job = job_1713870747829_0003

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/flight.db/flight_details/.hive-staging_hive_2024-04-23_05-52-43_707_603965056040236101-1/-ext-10000

Loading data to table flight.flight_details

Table flight.flight_details stats: [numFiles=1, numRows=1, totalSize=21, rawDataSize=20]

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 12.35 sec HDFS Read: 4324 HDFS Write: 98 SUCCESS

Total MapReduce CPU Time Spent: 12 seconds 350 msec

OK

Time taken: 732.464 seconds

hive> insert into flight_details values(102,"Delhi","Banglore",2024);

Query ID = cloudera_20240423060505_0739f610-4caf-4180-8268-c1be1e746513

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1713870747829_0004, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1713870747829_0004/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1713870747829_0004

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2024-04-23 06:12:41,862 Stage-1 map = 0%, reduce = 0%

2024-04-23 06:13:42,066 Stage-1 map = 0%, reduce = 0%

2024-04-23 06:14:42,918 Stage-1 map = 0%, reduce = 0%

2024-04-23 06:15:43,913 Stage-1 map = 0%, reduce = 0%

2024-04-23 06:16:44,719 Stage-1 map = 0%, reduce = 0%

2024-04-23 06:17:45,894 Stage-1 map = 0%, reduce = 0%

2024-04-23 06:18:47,797 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 5.69 sec

2024-04-23 06:19:25,313 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 12.13 sec

MapReduce Total cumulative CPU time: 12 seconds 130 msec

Ended Job = job_1713870747829_0004

Stage-4 is selected by condition resolver.

Stage-3 is filtered out by condition resolver.

Stage-5 is filtered out by condition resolver.

Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/flight.db/flight_details/.hive-staging_hive_2024-04-23_06-05-28_363_5397672270573654150-1/-ext-10000

Loading data to table flight.flight_details

Table flight.flight_details stats: [numFiles=2, numRows=2, totalSize=45, rawDataSize=43]

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 12.82 sec HDFS Read: 4328 HDFS Write: 101 SUCCESS

Total MapReduce CPU Time Spent: 12 seconds 820 msec

OK

Time taken: 859.643 seconds

hive> select * from flight_details;

OK

101	Pune	Mumbai	2022
-----	------	--------	------

102	Delhi	Banglore	2024
-----	-------	----------	------

Time taken: 1.019 seconds, Fetched: 2 row(s)

hive> select * from air_flight;

OK

101	20	Parth	Business class
-----	----	-------	----------------

102	21	Ayush	Economy class
-----	----	-------	---------------

Time taken: 0.632 seconds, Fetched: 2 row(s)

hive> describe flight_details;

OK

fno	int
-----	-----

source	varchar(20)
--------	-------------

destination	varchar(20)
-------------	-------------

year	int
------	-----

Time taken: 1.37 seconds, Fetched: 4 row(s)

hive> describe air_flight;

OK

fno	int
-----	-----

pno	int
-----	-----

pname	varchar(20)
-------	-------------

class	varchar(20)
-------	-------------

Time taken: 1.232 seconds, Fetched: 4 row(s)

```
hive> select a.fno, a.source, a.destination, a.year, b.pno, b.name, b.class from flight_details a join air_flight b
on (a.fno=b.fno);
```

FAILED: SemanticException [Error 10002]: Line 1:56 Invalid column reference 'name'

```
hive> select a.fno, a.source, a.destination, a.year, b.pno, b.pname, b.class from flight_details a join air_flight b
on (a.fno=b.fno);
```

Query ID = cloudera_20240423063232_389a1c6d-1119-457d-b747-dd1100fcb075

Total jobs = 1

Execution log at: /tmp/cloudera/cloudera_20240423063232_389a1c6d-1119-457d-b747-dd1100fcb075.log

2024-04-23 06:33:36 Starting to launch local task to process map join; maximum memory = 1013645312

2024-04-23 06:34:36 Dump the side-table for tag: 0 with group count: 2 into file:
file:/tmp/cloudera/3ea11894-0576-4ac6-be28-f0d5b008991b/hive_2024-04-23_06-32-26_489_3266736177048635829-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable

2024-04-23 06:34:37 Uploaded 1 File to: file:/tmp/cloudera/3ea11894-0576-4ac6-be28-f0d5b008991b/hive_2024-04-23_06-32-26_489_3266736177048635829-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile00--.hashtable (331 bytes)

2024-04-23 06:34:37 End of local task; Time Taken: 60.961 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1713870747829_0005, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1713870747829_0005/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1713870747829_0005

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2024-04-23 06:38:54,107 Stage-3 map = 0%, reduce = 0%

2024-04-23 06:39:54,329 Stage-3 map = 0%, reduce = 0%

2024-04-23 06:40:55,622 Stage-3 map = 0%, reduce = 0%

2024-04-23 06:41:58,273 Stage-3 map = 0%, reduce = 0%

2024-04-23 06:42:58,616 Stage-3 map = 0%, reduce = 0%

2024-04-23 06:43:58,915 Stage-3 map = 0%, reduce = 0%, Cumulative CPU 5.1 sec

2024-04-23 06:44:59,626 Stage-3 map = 0%, reduce = 0%, Cumulative CPU 5.1 sec

2024-04-23 06:45:13,465 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 11.73 sec

MapReduce Total cumulative CPU time: 11 seconds 730 msec

Ended Job = job_1713870747829_0005

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 12.14 sec HDFS Read: 7443 HDFS Write: 92 SUCCESS

Total MapReduce CPU Time Spent: 12 seconds 140 msec

OK

101	Pune	Mumbai	2022	20	Parth	Business class
102	Delhi	Banglore	2024	21	Ayush	Economy class

Time taken: 783.01 seconds, Fetched: 2 row(s)

```
hive> create index flight_index on table flight_info(flight_num) as  
'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler' WITH DEFERRED REBUILD;
```

OK

Time taken: 11.138 seconds

```
hive> show tables;
```

OK

air_flight

flight__flight_info_flight_index__

flight_details

flight_info

values__tmp__table__1

values__tmp__table__2

values__tmp__table__3

values__tmp__table__4

Time taken: 0.24 seconds, Fetched: 8 row(s)

```
hive> select day,month,year, AVG (dep_delay) as average_deparature_delay from flight_info where year=2008  
group by year, month, day;
```

Query ID = cloudera_20240423072222_2e536428-605b-4823-b452-8130de32d8c4

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks not specified. Estimated from input data size: 1

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job_1713870747829_0006, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1713870747829_0006/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1713870747829_0006

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

2024-04-23 07:26:50,937 Stage-1 map = 0%, reduce = 0%
2024-04-23 07:27:51,590 Stage-1 map = 0%, reduce = 0%
2024-04-23 07:28:52,360 Stage-1 map = 0%, reduce = 0%
2024-04-23 07:29:53,374 Stage-1 map = 0%, reduce = 0%
2024-04-23 07:30:54,171 Stage-1 map = 0%, reduce = 0%
2024-04-23 07:31:55,087 Stage-1 map = 0%, reduce = 0%
2024-04-23 07:32:56,092 Stage-1 map = 0%, reduce = 0%
2024-04-23 07:33:57,399 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 7.84 sec
2024-04-23 07:35:05,399 Stage-1 map = 0%, reduce = 0%, Cumulative CPU 7.84 sec
2024-04-23 07:35:33,324 Stage-1 map = 67%, reduce = 0%, Cumulative CPU 16.66 sec
2024-04-23 07:35:44,939 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 17.19 sec
2024-04-23 07:36:46,201 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 17.36 sec
2024-04-23 07:37:46,822 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 17.36 sec
2024-04-23 07:38:47,020 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 17.36 sec
2024-04-23 07:39:36,613 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 19.31 sec
2024-04-23 07:40:36,591 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 25.87 sec
MapReduce Total cumulative CPU time: 25 seconds 870 msec

Ended Job = job_1713870747829_0006

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 26.53 sec HDFS Read: 25126122 HDFS Write: 875
SUCCESS

Total MapReduce CPU Time Spent: 26 seconds 530 msec

OK

1	1	2008	17.57681842916742
2	1	2008	23.900056359195943
3	1	2008	19.370313695485844
4	1	2008	18.612678509230232
5	1	2008	25.976967114898148
6	1	2008	22.146653781106547
7	1	2008	14.395251396648044
8	1	2008	12.124760306807287
9	1	2008	5.839149336153214
10	1	2008	9.223829201101928
11	1	2008	9.410679275746743
12	1	2008	1.6842865395725015

13	1	2008	6.079343193782903
14	1	2008	4.633204633204633
15	1	2008	5.640961857379768
16	1	2008	1.9354166666666666
17	1	2008	18.21534910559723
18	1	2008	12.01187917185202
19	1	2008	7.5900463308922435
20	1	2008	6.213233458177278
21	1	2008	25.198426472289714
22	1	2008	17.538498383427136
23	1	2008	11.585463541053128
24	1	2008	9.975531671621313
25	1	2008	14.944508404328804
26	1	2008	4.631294964028777
27	1	2008	25.05219499744768
28	1	2008	14.486067019400354
29	1	2008	9.989655592065231
30	1	2008	6.108780661215784
31	1	2008	27.131638620360423

Time taken: 1109.855 seconds, Fetched: 31 row(s)

hive> select day,month,year,dest, AVG (dep_delay) as average_deparature_delay from flight_info where year=2008 and AVG (dep_delay)>0 group by year, month, day,dest;